Compiling Snort 3 Rule Headers into Conservative P4 Prefilters: An Evaluation on BMv2

Arthur de Freitas Medeiros
Universidade Federal do Pampa (UNIPAMPA)
Alegrete, Brazil
arthurfreitas.aluno@unipampa.edu.br

October 2026

Abstract

Programmable data planes can evaluate lightweight packet predicates before a software intrusion detection system (IDS), but Snort 3 rules also depend on payload, stream, flow/state, and other semantics that a P4 target cannot necessarily represent. We present a semantics-aware, target-aware Snort-to-P4 compilation approach that classifies each rule-to-predicate relation as Exact, Necessary, or Unsupported. A necessary predicate is meant to retain every packet satisfying its source rule within an explicitly enforced packet domain, while omitted semantics remain explicit in a manifest recording predicate-to-rule provenance. Packets outside that domain, such as those carrying IPv4 options, fragments, or truncated headers, are failed open as conservative candidates instead of being matched against fields the parser cannot locate. What a header condition guarantees depends on how the IDS evaluates it: once Snort 3 identifies a flow's service, it ignores the header ports and protocol of service-based rules, so the compiler keeps only their addresses. An earlier compilation that kept those conditions over-constrained 2,931 predicates, and one of its plans marked no-match, on BMv2, the witnesses of 190 HTTP rules on which Snort alerted off their header port. Under the corrected compilation, 3,996 of 4,017 rules are classified as necessary, 21 as unsupported, and none as exact, and backend lowering produces 235 P4 entries. The global binary prefilter marked 100% of parser-eligible packets as candidates on SQLi, XSS, and held-out Brute Force traces replayed on BMv2, and a strong-discriminator ablation that removes the broadest entries still marked 92.2–94.9%. Provenance supports an offline utility, per-packet candidate rule sets, which average about four of the 739 rules that ablation keeps P4-eligible; the 3,278 rules left in fallback apply to every packet, however, and the held-out reduction over the full ruleset is 18.3%, against 81.8% under the earlier compilation. Observed alerts were preserved in the traces, and in an end-to-end campaign over 1,000 sampled rules, with Snort 3 as the oracle and BMv2 as the target, no rule that a plan prefilters alerted on a witness the data plane marked no-match. These

results show that semantic safety and operational selectivity are distinct, and that much of the selectivity a header-derived prefilter appears to offer can come from conditions the IDS does not check; candidate-set-aware IDS execution remains future work.

1 Introduction

Signature-based network IDSs remain dependent on rules that combine packet headers with payload inspection, stream and flow state, service semantics, and temporal conditions. These conditions cannot generally all be represented by a programmable data-plane target. The resulting challenge is not merely to translate rule syntax to P4, but to partially compile Snort 3 detection semantics without obscuring what the target cannot represent.

Programmable data planes can evaluate lightweight packet predicates before software IDS processing. Prior P4-based IDS and prefiltering efforts motivate this form of data-plane assistance [7, 14], while work on rule-based prefiltering emphasizes the importance of retaining the context required by the IDS [3]. However, the interaction between semantic preservation, target limitations, and ruleset-scale selectivity remains insufficiently characterized.

A Snort rule may contain representable header predicates together with payload, stream, flow/state, normalization, or service-dependent conditions. Compiling only the representable subset without recording the omitted semantics risks presenting a partial implementation as equivalent to the original rule. A safe necessary prefilter instead establishes, for a rule r and generated predicate p_r, the relation $r \Rightarrow p_r$ under a stated semantic profile and a stated packet domain; it does not assume $p_r \Rightarrow r$ or $r \Leftrightarrow p_r$. Traffic the data plane cannot parse as intended is retained by failing open rather than being matched against fields read from the wrong offset. The detailed model is developed in the problem formulation.

Semantic safety alone, however, does not determine operational usefulness. When thousands of necessary predicates are aggregated into a global binary candidate decision, their union can cover nearly all traffic in a workload. Semantic coverage and packet-level selectivity must therefore be evaluated as separate properties. Moreover, the binary decision discards which predicate and which source rule produced a match. Retaining predicate provenance enables an alternative question: not only whether a packet remains a candidate for IDS handling, but which IDS rules remain plausible candidates for that packet. This is evaluated here as an offline candidate-rule-space opportunity, not as an implemented runtime execution mechanism.

This paper develops and evaluates a semantics-aware, target-aware compilation workflow from Snort 3 rules to P4 predicates, with explicit delegation and fallback. Under the evaluated configuration, 3,996 of 4,017 rules admitted necessary predicates. For the 2,931 service-based rules among them, those predicates keep only the header addresses, because Snort 3 stops checking ports and protocol once it identifies the flow's service; an earlier compilation that kept

them is reported for what it missed. The global binary prefilter marked 100% of parser-eligible packets as candidates in the three evaluated traces. Provenance-based analysis reduced the mean per-packet candidate set by more than 99% within the 739 rules that a strong-discriminator ablation keeps P4-eligible, but by only 18.29–18.30% over the full ruleset; under the earlier compilation the held-out reduction over the full ruleset was 81.81%. These observations motivate the following contributions and research questions.

1.1 Contributions

This work makes the following contributions:

- A semantics-aware, target-aware method for deriving P4-executable predicates from Snort 3 rules and classifying their relation as EXACT, NECESSARY, or UNSUPPORTED.
- A compilation manifest that records preserved and delegated predicates, fallback and unsupported rules, and provenance linking final P4 entries to source Snort rule identifiers.
- A P4 backend that performs predicate expansion, normalization, target-aware lowering, deduplication, and P4Runtime installation-plan generation; the evaluated backend reduces 4,040 predicate instances to 235 final installation entries as installation state.
- An explicitly enforced packet domain that gates predicate evaluation on IPv4 header form and fails open for options, fragments, truncated headers, and parser errors, validated by a controlled data-plane suite and by trace-level replay accounting.
- An empirical evaluation that separates semantic safety from ruleset-scale packet-level selectivity, including a causal analysis of broad predicates and distributed union coverage.
- A provenance-based, offline candidate-rule-space analysis, which reduces the mean per-packet candidate set by more than 99% within the P4-eligible partition but by at most 18.40% over the full ruleset, and a measurement of the candidate set that a grouped prefilter would forward.
- An analysis, from the Snort 3.12.2.0 source, showing that service-based rules, whether they declare a `service` option or acquire a service through their options, do not make their header ports and protocol necessary conditions; a compilation that follows this rule for the 2,931 affected rules; and a measurement, on BMv2 and over the traces, of what an earlier compilation that kept those conditions missed and of the selectivity it had appeared to offer.
- Empirical validation of semantic-split behavior and observed P4 alert-packet preservation, without claiming formal equivalence or dynamic stateful candidate-set execution.

- An end-to-end differential campaign that synthesizes a witness flow per rule and compares Snort 3 alerts with BMv2 verdicts frame by frame, observing no violation over a stratified sample of 1,000 rules under either plan and measuring how much the evaluated plans can falsify.

1.2 Research Questions

RQ1 – Semantic Coverage. To what extent can Snort 3 rules be safely represented by predicates executable in a programmable P4 data plane while explicitly preserving semantic limitations?

RQ2 – Backend Feasibility. Can the semantically derived predicates be lowered and installed in a practical P4 target without hiding unsupported semantics or exceeding target constraints?

RQ3 – Operational Selectivity. How selective is a binary P4 prefilter when thousands of semantically safe predicates are jointly deployed?

RQ4 – Detection Preservation and Alternative Utility. When packet-level bypass selectivity is limited, can the compiled predicates still preserve observed IDS alerts and reduce the per-packet candidate rule space?

1.3 Paper Organization

The remainder of the paper presents background and related work, followed by the problem formulation and proposed approach. It then describes the experimental methodology, results, discussion, threats to validity and limitations, conclusion, and the experimental-profile appendix.

2 Background

2.1 Snort 3 Rule Semantics

Snort [11] is a signature-based network IDS; this work targets Snort 3 [17]. A Snort 3 text rule has a header and a body. The header names an action, a protocol, the source address and port, a direction that is either unidirectional or bidirectional (`<>`), and the destination address and port; addresses and ports may be variables such as `$HOME_NET` or the wildcard `any`. The body, in parentheses, holds options: payload options such as `content` and `pcre` that inspect a buffer, flow options such as `flow` and `flowbits` that depend on session state, and metadata such as `sid`. Snort selects candidate rules with a fast-pattern search and then evaluates each candidate's options left to right, stopping at the first one that fails.

Snort 3 also distinguishes port-based from service-based rules. A rule is service-based when it declares a `service` option, uses an option specific to a service, or names a service protocol in its header, as in `alert http`. Every header component of a port-based rule must match. Once Snort has identified the flow's service as one of the services of a service-based rule, it checks the rule's header addresses but not its ports or protocol; a rule that names the service

protocol in its header keeps its port checks (Section 4.5). A header predicate taken from a rule is therefore an approximation whose relation to the rule has to be stated rather than assumed.

2.2 P4 Programmable Data Planes

P4 [2] programs describe how a data plane parses packets and processes them through match-action tables; this work uses $P4_{16}$ [15]. A parser extracts headers into typed fields, and control blocks apply tables whose keys use match kinds such as `exact`, `ternary`, and `lpm`; the v1model architecture adds `range` and `optional`. The program fixes which fields a table can match, while the control plane installs the entries at run time. The evaluated target is BMv2 [10], whose `simple_switch_grpc` process implements the v1model architecture and serves the P4Runtime API [16]. Through it a controller installs the pipeline configuration and writes table entries, and ternary entries carry a priority that orders overlapping matches. The evaluated backend uses target-supported Ethernet/IPv4 header fields as a candidate marker; it does not replace Snort detection.

2.3 Semantic Preservation and Prefiltering

Necessary predicates can safely retain traffic for Snort inspection, but their joint deployment need not provide useful packet-bypass selectivity.

3 Related Work

3.1 Rule-Derived and Hybrid P4 IDS Offloading

Rule-derived filtering has established that programmable data planes can assist software IDS processing. P4ID [7] derives packet filters from Snort rules using information such as IP addresses, transport-layer ports, protocols, and TCP flags, and combines stateless and stateful packet-processing stages. It therefore demonstrates both rule-derived filtering and stateful data-plane assistance. P4-ONIDS [14] addresses constrained data planes through range representation, rule aggregation, deduplication, and flow information to limit packets forwarded to the NIDS.

4MIDable [8] provides a framework for integrating existing security middleboxes with P4/SDN infrastructure. Its P4ID agent supports stateful processing and feedback to an unmodified IDS, and was evaluated on both software BMv2 and hardware Tofino targets. The work reports traffic-processing reductions without affecting observed detection rates. P4NIPS [12] provides a recent P4 network-intrusion prevention implementation. Together, these systems show the feasibility of rule-derived filtering, hybrid IDS/data-plane integration, stateful assistance, and resource-aware P4 representations; they are not treated here as evidence for a common per-rule semantic relation between a source IDS rule and a generated predicate.

3.2 Context Preservation and Hybrid IDS Processing

Brum et al. [3] revisit rule-based NIDS pre-filtering by considering traffic required for subsequent IDS processing even when a packet does not directly satisfy extracted rule conditions. Their eRBF approach incorporates flow context, including TCP connection establishment and termination information and selected protocol dependencies. Evaluated with Snort 3 and large rulesets, it shows why aggressive pre-filtering must account for context required by the IDS. Its software trace-based evaluation, however, does not address the mapping of those semantics to a P4 target's processing and resource constraints.

RNA [5] follows a complementary event-oriented direction: it analyzes selected Zeek-script events and maps them to protocol and event templates for programmable forwarding devices, together with complementary host-side mechanisms. This coordinated generation illustrates how selected IDS functionality can be distributed across host and data-plane components. Its scope is selected Zeek events and predefined templates rather than a general semantic classification of signature-based IDS rules according to the guarantees preserved by their data-plane representation.

3.3 Programmable-Data-Plane Security Landscape

The broader programmable-data-plane literature situates IDS offloading among other security applications and target models. Surveys of P4 fundamentals and applications [6], P4 security applications and their vulnerability context [1], and programmable-data-plane IDS approaches [4] document the breadth of the field. Implementation-oriented work also emphasizes that mapping security functionality to programmable targets involves resource and implementation constraints [9]. These sources motivate target-aware treatment of security functionality; they do not by themselves establish properties of the individual systems compared above.

3.4 Positioning of This Work

Prior systems demonstrate that rule-derived filtering, stateful assistance, resource-aware P4 representations, and coordinated IDS/data-plane execution are feasible. This work focuses on making the semantic consequences of partial compilation explicit. It classifies the relation between each source Snort rule and generated P4 predicate, records preservation, delegation, and fallback in a compilation manifest, and retains source-rule provenance through backend lowering and deduplication.

The evaluation also separates semantic safety from aggregate operational selectivity: a safe necessary predicate can be retained for Snort inspection yet contribute to a global binary candidate union that is not selective. Rather than presenting that outcome as a successful traffic-bypass result, the analysis uses retained provenance to examine an offline per-packet candidate-rule-space opportunity. Thus, the contribution is not Snort-to-P4 translation alone, but explicit

semantic and target-aware characterization of what is preserved, delegated, and still operationally useful. Table 1 compares these dimensions across the reviewed work.

Table 1: Related-work comparison. "Not explicit" denotes information not established in the reviewed scope, rather than absence of the feature.

Work	**Source IDS / rules**	**P4 target**	**Rule-derived filtering**	**State / context handling**	**Explicit semantic relation**	**Provenance retained**	**Main objective**
P4ID	Snort rules	P4 data plane	Yes	Yes (stateful stage)	Not explicit in reviewed scope	Not explicit in reviewed scope	Rule-derived IDS filtering
P4-ONIDS	NIDS rules	P4 data plane	Yes	Partial (flow information)	Not explicit in reviewed scope	Not explicit in reviewed scope	Resource-aware NIDS pre-filtering
4MIDable	Existing security middleboxes / P4ID agent	BMv2; Tofino	Yes	Yes	Not explicit in reviewed scope	Not explicit in reviewed scope	Hybrid security-VNF offloading
eRBF	Snort 3 rules	N/A (software trace model)	Yes	Yes	Not explicit in reviewed scope	Not explicit in reviewed scope	Context-preserving NIDS pre-filtering
RNA	Zeek events	P4 device + host	N/A	Partial (event / host coordination)	Not explicit in reviewed scope	Not explicit in reviewed scope	Event offload
P4NIPS	Not explicit in reviewed scope	P4	Not explicit in reviewed scope	Not explicit in reviewed scope	Not explicit in reviewed scope	Not explicit in reviewed scope	NIPS implementation
This work	Snort 3 rules	BMv2	Yes	Delegated to Snort	Yes (three classes)	Yes	Semantic transparency and selectivity analysis

4 Problem Formulation

4.1 Snort Rule Model

A signature-based IDS rule can combine multiple conditions that depend on different types of information. Some conditions can be evaluated directly from packet headers, while others may depend on payload contents, flow state, temporal relationships, protocol normalization, or information collected across multiple packets. Therefore, a rule cannot generally be treated as a single match-action entry when considering its execution in a programmable data plane.

For the purposes of this work, a rule r is represented as:

$$r = (H_r, P_r, F_r, T_r, N_r, a_r, m_r),$$

where H_r represents header predicates, P_r payload-dependent predicates, F_r flow and state conditions, T_r temporal conditions, N_r normalization or protocol-dependent conditions, a_r the associated action, and m_r rule metadata such as its identifier and priority. Given a target data plane s, only a subset of the predicates that compose a rule may be executable by that target. We denote by C_s the set of capabilities exposed by the target, such as exact, ternary, or range matching over selected packet-header fields. The compilation process analyzes each rule predicate with respect to C_s and produces a data-plane predicate p_r together with a semantic classification.

4.2 P4 Predicate Model and Semantic Relations

In this initial model, the generated representation is assigned to one of three classes:

- **EXACT**: all predicates considered by the supported rule profile are represented without semantic relaxation in the generated data-plane predicate.
- **NECESSARY**: the generated predicate preserves only conditions that are necessary for the original rule to match. Therefore, packets that satisfy the generated predicate still require inspection by the original IDS.
- **UNSUPPORTED**: the compiler cannot derive a safe data-plane predicate under the capabilities assumed for the target, and evaluation remains delegated to the IDS.

4.3 Safety Condition

The semantic classification determines how the generated component can be interpreted. An EXACT representation records no delegated, fallback, or unassessed detection predicate within the supported profile. The current P4 program nevertheless implements candidate marking through the single `send_to_ids` action rather than a general data-plane detection action. A NECESSARY representation is used only as a conservative pre-filter: matching traffic remains subject

to inspection by Snort, since the data-plane predicate alone is insufficient to reproduce the complete detection decision. An UNSUPPORTED rule remains delegated to Snort. Therefore, the existence of a data-plane match is not treated as equivalent to an IDS alert unless the required semantic guarantees are satisfied. For a rule r and a generated NECESSARY predicate p_r, the safety relation is $r \Rightarrow p_r$; the converse $p_r \Rightarrow r$ is not assumed. UNSUPPORTED means that no safe selective P4 predicate was established under the target profile, not that the rule is ignored.

4.4 Predicate-Level Soundness of Partial Compilation

The NECESSARY claim is a predicate-level statement under the current Ethernet/IPv4 target profile, not a general claim about arbitrary Snort execution. For a parsed rule r, let H_r denote the conjunction of its header conditions that the capability analysis marks PRESERVED or EXPANDED; let U_r denote the remaining detection-relevant conditions recorded as delegated, fallback, or unassessed. A source-rule match entails every header condition that Snort checks for that rule, and the compiler restricts H_r to those conditions, so that $r(x) \Rightarrow H_r(x)$; Section 4.5 identifies them for the rules whose header Snort does not evaluate in full. The backend materializes a predicate $p_r = L(H_r)$ only for a prefilter candidate. Under the implementation assumptions that the accepted header forms are parsed and normalized as recorded in the manifest, and that lowering does not exclude a packet satisfying H_r, the construction gives

$$r(x) \Rightarrow H_r(x) \Rightarrow L(H_r)(x) = p_r(x).$$

The implication uses no claim about the omitted U_r conditions: their presence can make p_r broader, which is why NECESSARY is not EXACT. If the compiler cannot establish a selective H_r, it emits no prefilter entry and classifies the rule as UNSUPPORTED. EXACT would additionally require that all detection-relevant predicates be represented without relaxation; no such rule occurred in the evaluated Community ruleset.

4.5 Service-Based Rules

The premise $r(x) \Rightarrow H_r(x)$ holds when Snort evaluates the whole rule header, and Snort 3 does not always do so. In Snort 3.12.2.0, before the header of a rule with services is checked, Snort compares the service it has identified for the flow with the rule's services. If the identified service is one of them, the port checks are skipped and the protocol check succeeds unconditionally, while the address checks still run; if a different service has been identified, the rule is rejected; only when no service has been identified is the whole header checked (`src/detection/detection_options.cc` and `src/detection/rtn_checks.cc` in the Snort 3.12.2.0 source). Snort's documentation states a coarser form of this rule: for a service-based rule, evaluation of the header is skipped [17]. Service identification can come from content inspection, so an HTTP request on a port

outside `HTTP_PORTS` is still evaluated against the HTTP rules. One form is an exception. A rule that names a service protocol in its header, such as the twelve rules of the form `alert http ... -> ...` in the evaluated ruleset, is marked as a user-mode rule when Snort builds its header checks, and that restores the port checks (`src/parser/parse_rule.cc` and `src/detection/fp_detect.cc`).

A rule need not declare its services. When the rules are loaded, an option that selects a buffer registered by a service inspector adds that inspector's services to a rule that has none, and a few option-name prefixes add a guessed service (`src/detection/fp_utils.cc`). The DCE options `dce_iface` and `dce_stub_data`, for instance, add `dcerpc` and `netbios-ssn`.

For a service-based rule whose header names an IP protocol, the header condition that is necessary for an alert is therefore H_r^{addr}, the address part of its header; the full header is necessary only for packets whose service Snort has not identified. The compiler therefore represents these rules by H_r^{addr} and delegates their ports and protocol to Snort. It keeps the protocol only when the rule's `flags` option already restricts it to TCP, since that option matches no other protocol. From here on, *service-based* refers to these rules only; the user-mode rules keep their port checks and are compiled, like port-based rules, with their full header. An earlier version of the compiler represented every rule by its full header; Section 7.2 measures how many rules that affected and what it changed.

4.6 Enforced Packet Domain

The implication above presupposes that the data plane reads the fields the predicate refers to. That presupposition is not free: a parser that assumed a fixed 20-byte IPv4 header would read IPv4 option bytes as transport ports, so a packet satisfying r could fail p_r and the necessary relation would not hold. The implementation therefore makes the precondition explicit rather than assumed.

Let D denote the *supported packet domain*: untagged Ethernet carrying direct IPv4 with version = 4, IHL = 5, fragOffset = 0, flags either zero or `DF`, and no parser error. The P4 parser gates L4 extraction on membership in D, and the ingress control re-checks the same conditions before applying the predicate table. For $x \notin D$, the pipeline executes `packet_domain_fallback`, which marks the packet as a candidate unconditionally and counts it on a dedicated index. Writing $\text{cand}(x)$ for the candidate outcome of the deployed pipeline:

$$r(x) \Rightarrow \begin{cases} p_r(x) \Rightarrow \text{cand}(x), & x \in D, \\ \text{cand}(x), & x \notin D. \end{cases}$$

Consequently, no packet satisfying the rule is excluded from Snort inspection, whether or not the parser can interpret it. The data plane fails open by construction: a no-match verdict is emitted only inside D, because no-match is the only outcome that could ever motivate bypassing the IDS. This makes the domain a stated precondition of the selectivity metrics rather than a hidden assumption, and it is why candidate rates in Section 7 are reported over D with domain fallback counted separately.

4.7 Semantic Safety versus Operational Selectivity

A predicate can be semantically safe as a necessary condition while still matching most packets once it is jointly deployed with thousands of other predicates. This distinction is evaluated explicitly in the results.

5 Proposed Approach

The proposed approach introduces a semantics-aware compilation pipeline between Snort 3 rules and a programmable data plane. Instead of directly converting a complete rule into P4 code, the pipeline first parses and decomposes the rule into an intermediate representation. Each predicate is then analyzed according to the capabilities of the selected P4 target. The result of this analysis determines which components can be represented in the data plane and which must remain delegated to Snort.

The proposed pipeline consists of five main stages: (i) Snort rule parsing, (ii) construction of an intermediate representation, (iii) target capability analysis, (iv) semantic classification, and (v) generation of the corresponding data-plane configuration or IDS fallback. Figure 1 places these stages between their inputs and the BMv2 target.

5.1 Snort Rule Parsing and Intermediate Representation

Stages (i) and (ii) receive Snort 3 rules and extract the elements required for subsequent semantic analysis. Rather than translating the textual rule directly into P4 constructs, the parser converts each rule into a structured intermediate representation (IR). This representation separates the original IDS syntax from the capabilities of a particular data-plane target, allowing the same rule to be analyzed independently of its final implementation.

For the initial prototype, the parser focuses on a restricted rule profile containing the rule action, network protocol, source and destination addresses, source and destination ports, traffic direction, TCP flags, and rule identifier. Additional rule options are also identified when possible, even when they are not executable by the initial P4 backend. This allows unsupported predicates to remain visible during semantic analysis instead of being silently removed from the original rule. Conceptually, the IR of a rule contains the following information:

```
rule:
  sid: 1001
  action: alert
  protocol: tcp
  source:
    address: any
    port: any
  direction: ->
```

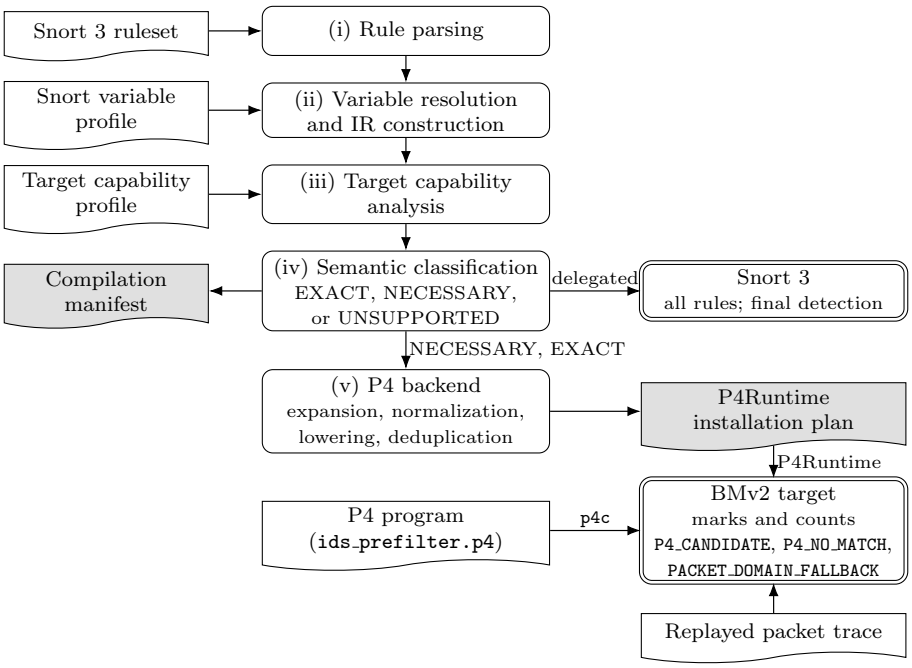


Figure 1: Implemented compilation pipeline and its artifacts. Rounded boxes are the compiler stages (i)–(v); wavy-bottom shapes are inputs (white) and generated artifacts (gray); double-bordered boxes are execution targets. All rules, including UNSUPPORTED ones, remain in Snort 3, which performs final detection. The deployed P4 program marks candidate metadata and counts each packet in one of three categories; it is not integrated into a forwarding or bypass path in front of Snort, and packet traces are replayed into BMv2 for measurement.

```
destination:
  address: 10.0.0.2
  port: 80
predicates:
  - type: content
    value: "attack"
```

By preserving both supported and unsupported rule components in the IR, the subsequent stages can explicitly determine which semantics are retained by the generated data-plane predicate. The IR therefore acts as an interface between IDS-specific rule parsing and target-specific code generation.

5.2 Target Capability Analysis

Stage (iii) evaluates each predicate represented in the IR against a predefined capability profile of the target data plane. This profile describes which operations and fields are considered executable by the initial P4 backend. Separating target capabilities from the rule parser prevents the frontend from assuming that every parsed IDS construct can be implemented by the programmable switch.

For the initial prototype, the target profile is intentionally restricted to predicates that can be evaluated using packet-header information. The supported subset includes IPv4 addresses and subnets, TCP and UDP ports, IP protocol information, traffic direction, and TCP flags. Predicates requiring general payload inspection, regular expressions, stream reassembly, arbitrary flow state, or protocol-specific normalization are not considered executable by the initial backend. A predicate marked as unsupported by this profile is not removed

Table 2: Initial capability profile considered by the prototype.

Rule component	**Initial support**
IPv4 protocol	Supported
Source/destination IP	Supported
IP subnet	Supported
TCP/UDP ports	Supported
Traffic direction	Supported
TCP flags	Supported
Payload content	Unsupported
PCRE	Unsupported
Stream reassembly	Unsupported
Arbitrary flow state	Unsupported

from the IR. Instead, its presence influences the semantic classification of the generated representation. This distinction allows the compiler to use supported predicates for pre-filtering while retaining an explicit record of the conditions that must still be evaluated by Snort. Table 3 lists, construct by construct, how the implemented Ethernet/IPv4 profile compiles each form and what it guarantees.

Table 3: Compilation semantics in the implemented Ethernet/IPv4 profile. "Omitted" means absent from p_r while remaining explicit in the manifest and delegated to Snort; it does not mean ignored by the IDS.

Snort construct	**Compiler handling**	**P4 representation**	**Treatment**	**Guarantee / condition**
`tcp`, `udp`, `icmp`, `ip`	Protocol name is assessed; `ip` is a wildcard L4 protocol.	Exact IPv4 protocol ternary key, or no L4 key for `ip`.	Preserved	Within the Ethernet/IPv4 parser scope.
IPv4 address or CIDR; resolved network variable	Resolver substitutes whole-field variables; CIDR is canonicalized to network and prefix length.	IPv4 ternary value/mask.	Preserved	CIDR membership is retained; unsupported or negated expressions do not reach this form.
Exact port or finite port/address list	Resolver flattens lists; finite alternatives are expanded by Cartesian product.	One exact ternary key per concrete alternative.	Preserved / expanded	The union of concrete entries retains the finite-list alternatives, for packets inside the enforced domain D of Section 4.6.
`->` or `<>`	Forward orientation is retained; bidirectional rules swap source/destination addresses, ports, and expansion fields for a reverse entry.	Forward entry; plus reverse entry for `<>`.	Preserved / expanded	Both directions are materialized for `<>`; `<-` falls back.
Supported TCP flags	Flag syntax is normalized to masks; `*` is lowered to a union of ternary entries.	TCP-flag ternary value/mask.	Preserved / expanded	Valid only for forms accepted by the flag normalizer, with TCP protocol validation, and for packets inside D.
Negation, port ranges, unresolved complex header forms	Capability analysis records fallback or unassessed treatment.	Omitted from p_r.	Fallback / delegated	Omission can broaden a predicate; absent another selective H_r, the rule is UNSUPPORTED.
Payload, PCRE, buffers, flow/state, stream, temporal, and application semantics	Ordered options and buffer context remain in the IR and manifest.	No P4 match.	Delegated	Snort remains responsible for these conditions and for final detection.
Service-style or otherwise nonselective rule	No traditional header is invented; a rule without a selective represented header is not emitted.	No P4 entry.	UNSUP.	Entire rule remains in the Snort fallback.
Traditional header on a service-based rule	Addresses assessed as for a port-based rule; ports and protocol delegated, except TCP under `flags`.	Address keys, plus TCP in that case.	Preserved / delegated	Necessary also once Snort identifies the flow's service (Section 4.5).
IPv4 options, fragments, truncated headers, non-IPv4 frames	Outside the compiler's scope; handled by the deployed program rather than by lowering.	No table lookup; `packet_domain_fallback` marks a candidate on a dedicated counter.	Fail-open	Retained for Snort without a no-match verdict, so misparsed L4 fields cannot break $r \Rightarrow$ cand.

5.3 Semantic Classification and Decision

After capability analysis, the compiler determines the semantic class of the generated data-plane representation. The decision considers both the predicates that can be executed by the target and those that remain outside its capability profile. The objective is to prevent a partial translation from being interpreted as an implementation of the complete IDS rule.

If all predicates required by a rule within the considered rule profile can be represented without relaxation, the generated representation is classified as EXACT. If only a subset of necessary predicates can be represented, while additional conditions must still be evaluated by Snort, the representation is classified as NECESSARY. Finally, when no safe and useful predicate can be derived under the target profile, the rule is classified as UNSUPPORTED and remains delegated to the IDS. For example, consider the following simplified Snort rule:

```
alert tcp any any -> 10.0.0.2 80
    (content:"attack"; sid:1001;)
```

The initial target profile can represent the protocol, destination address, and destination port, but does not support the content predicate. Therefore, the generated data-plane predicate corresponds to:

$$p_r = (\text{protocol} = \text{TCP}) \wedge (\text{dstIP} = 10.0.0.2) \wedge (\text{dstPort} = 80).$$

Since matching p_r is not sufficient to establish that the payload contains the string required by the original rule, the generated representation is classified as NECESSARY. Matching packets are therefore candidates for further inspection by Snort rather than detections produced directly by the data plane. The classification result is recorded together with the predicates preserved in the data plane and those delegated to the IDS. This information forms a per-rule compilation manifest, providing an explicit description of the semantic decisions made during translation.

5.4 Compilation Manifest

For each processed rule, the compiler produces a compilation manifest that records the result of the semantic analysis. The manifest identifies the original rule, its assigned semantic class, the predicates preserved by the data-plane representation, and the predicates that remain delegated to Snort. It also records the intended use of the generated component, providing an auditable description of the transformation performed by the compiler.

For the example rule presented previously, a simplified compilation manifest can be represented as follows:

```
rule_id: 1001
classification: NECESSARY
```

```
preserved:
  - protocol: tcp
  - dst_ip: 10.0.0.2
  - dst_port: 80

delegated:
  - content: "attack"

data_plane_action: PREFILTER
final_detection: SNORT
```

The manifest prevents unsupported predicates from being silently discarded during compilation. In addition to supporting debugging and inspection of the generated configuration, it provides information that can later be used to validate whether the behavior observed during execution is consistent with the semantic classification assigned by the compiler.

5.5 P4 Backend, Lowering, and Runtime Installation

The backend lowers target-supported predicates to P4 table matches, deduplicates identical install state, and emits a P4Runtime installation plan. Install-entry provenance retains the associated source-rule IDs for later candidate rule set analysis.

For a finite list, the expansion stage produces the Cartesian product of the approved alternatives; for `<>`, it materializes both source/destination orientations before that expansion. Normalization converts IPv4 CIDRs to network/prefix objects, ports to exact integers, and accepted TCP-flag syntax to masks. TCP-flag lowering retains exact, all-set, all-clear, and at-least-one semantics through ternary matches, with the last represented as a union of entries. P4Runtime encoding then checks field widths and encodes CIDR prefixes as ternary masks.

Deduplication merges only identical normalized or lowered matches, all of which execute the same positive `send_to_ids` action. It therefore does not change the logical union of packets marked candidate, although it does reduce installation state. When identical matches originate from several rules, their `source_rule_ids` are unioned in the retained provenance; otherwise later per-packet candidate-rule-set analysis would lose rule-level information.

The deployed P4 program enforces the packet domain D of Section 4.6 around this table. The parser transitions to L4 extraction only for direct IPv4 with IHL = 5 and no fragmentation, and the ingress control repeats the test, additionally rejecting parser errors, before the table is applied. Three counter indices are kept distinct: `P4_CANDIDATE` for a table hit inside D, `P4_NO_MATCH` for a miss inside D, and `PACKET_DOMAIN_FALLBACK` for traffic outside D, which is marked candidate unconditionally. Collapsing the third category into the second would overstate selectivity, and collapsing it into the first would hide the parser's actual coverage; the metrics below therefore report all three.

5.6 Candidate Rule Provenance and Prototype Scope

The initial prototype deliberately restricts the compilation backend to the capability profile described in Table 2. The objective of this first implementation is not to reproduce the complete Snort 3 rule language in P4, but to evaluate whether explicit semantic classification can support safe rule-derived pre-filtering.

More complex constructs, including general payload inspection, PCRE, stream reassembly, arbitrary stateful dependencies, and protocol-specific normalization, remain outside the initial backend and are delegated to Snort. Future extensions may introduce additional target profiles and backends without requiring changes to the rule representation and semantic-analysis model.

6 Experimental Methodology

6.1 Research Questions and Evaluation Design

The evaluation is designed to investigate two main aspects of the proposed approach: whether the compiler records the intended predicate-level semantic relation between Snort 3 rules and their generated data-plane predicates, and whether the resulting pre-filtering configuration preserves observed alerts within the evaluated scope.

The methodology is divided into three stages. First, a set of Snort 3 rules is processed by the proposed compiler to measure the coverage of the initial target profile and the distribution of semantic classifications. Second, the generated predicates are executed using a P4 software target. Finally, the behavior of the resulting pre-filtering pipeline is compared with the execution of the original rules directly in Snort, both on captured traces and in an end-to-end campaign over synthesized flows (Section 6.5).

6.2 Experimental Environment and P4/BMv2 Target

The initial experimental environment consists of Snort 3 as the reference IDS and a P4 software switch as the programmable data-plane target. The same traffic traces and relevant detection rules are used in both the baseline and the proposed pipeline. The software-based environment is selected to enable reproducible validation of the compilation process before considering hardware-specific optimizations.

The frozen results were produced with Snort 3 version 3.12.2.0, the P4 compiler `p4c` version 1.2.4.2 (build SHA `624c6be80`) invoked through its `p4c-bm2-ss` backend, and BMv2 `simple_switch_grpc` version 1.15.0 (build `68f4a978`), on Ubuntu 22.04 LTS with Python 3.10. The P4Runtime endpoint is `127.0.0.1:9559` with device identifier 0. The data-plane interfaces are `veth` pairs configured with an MTU of 65535, because the replay path writes frames larger than the default MTU, and with IPv6 disabled on each interface, so that router solicitation and multicast traffic generated by the host stack is not

counted by the data-plane counters. The script `scripts/setup_bmv2.sh` in the artifact repository reproduces this configuration.

Two execution paths are compared:

1. **Baseline**: the original traffic is processed directly by Snort using the original rules.
2. **Proposed pipeline**, evaluated by composition rather than as one deployed system: BMv2 classifies traffic with the installed predicates, and Snort processes the same traffic separately with the semantic candidate partition and with the UNSUPPORTED partition. The preservation measures defined below relate the two, by comparing the union of the partition runs with the baseline and by checking the P4 classification of the packets on which candidate alerts fell. No component forwards traffic from BMv2 to Snort.

6.3 Ruleset, Profiles, Traces, and Metrics

The frozen experimental Snort variable profile is documented for reproducibility in Appendix A.

Every result below is for the compilation of Section 4.5, which keeps only the address part of a service-based rule's header. An earlier compilation kept ports and protocol for those rules. Its artifacts are archived and reproducible byte for byte, and its results are reported only where they show what the correction changed (Section 7.2).

The evaluation considers four main metrics. First, *rule coverage* measures the proportion of analyzed Snort rules for which the compiler can generate an EXACT or NECESSARY data-plane representation. Second, *semantic classification distribution* reports the proportion of rules assigned to each class: EXACT, NECESSARY, and UNSUPPORTED.

Third, *alert preservation* compares the alerts produced by the proposed pipeline with those produced by the baseline Snort execution. For the evaluated traffic, alert preservation is defined as:

$$AP = \frac{|A_{\text{pipeline}} \cap A_{\text{baseline}}|}{|A_{\text{baseline}}|},$$

where A_{baseline} is the set of alerts generated by Snort over the original traffic and A_{pipeline} is the union of the alerts from the two partition runs of the proposed pipeline, the semantic candidate partition and the UNSUPPORTED partition, over the same traffic.

Finally, *candidate traffic* records parser-eligible packets marked candidate or no-match by the P4 prefilter. It is reported as a selectivity measurement, not as measured runtime reduction:

$$CS = \frac{N_{\text{candidate}}}{N_{\text{parser-eligible}}},$$

where $N_{\text{candidate}}$ is the number marked candidate and $N_{\text{parser-eligible}}$ is the parser-eligible population. Here, a parser-eligible packet is one inside the enforced

domain D of Section 4.6, that is, one for which the implemented P4 parser produces the Ethernet/IPv4 headers required by the predicate evaluation. A candidate packet matches at least one installed predicate, whereas a no-match packet is parser eligible but matches none. Candidate and no-match rates use only parser-eligible packets as their denominator. Packets outside D are counted as `PACKET_DOMAIN_FALLBACK`: they are retained as candidates for Snort but are excluded from this denominator, since the prefilter never claims selectivity over traffic it cannot parse.

Packet-domain validation. Because the metric above depends on the parser and the offline matcher agreeing, both are validated directly. A controlled data-plane suite sends one crafted frame per case through BMv2 with the 235-entry global plan installed and compares the resulting counter delta with the offline expectation. The eleven cases cover TCP, UDP, and ICMP candidates, a guaranteed no-match, exact and masked TCP flags, IPv4 options (IHL $\neq 5$), a first fragment with MF set, a non-initial fragment, a truncated IPv4 header, and a truncated TCP header. Separately, each full trace is replayed through BMv2 over a raw AF_PACKET socket that transmits the captured Ethernet bytes without reserialization, with a 2 ms inter-frame delay and veth transmit queues of 10,000 frames. Each replay starts from a fresh pipeline installation, with counters verified at zero and an idle audit confirming a zero delta; sender frame counts are then reconciled against the sum of the three counter categories.

Installation-state reduction. The backend artifact distinguishes `source_entries`, namely the `generated/expanded/normalized` predicate instances before installation-state reduction, from `final_install_entries`, namely the final P4 table entries after deduplication and TCP-flag lowering. Let these quantities be N_{source} and N_{install}, respectively. We measure

$$ISR = \left(1 - \frac{N_{\text{install}}}{N_{\text{source}}}\right) \times 100\%.$$

The backend reports $N_{\text{source}} = 4040$ and $N_{\text{install}} = 235$, which yields a 94.18% reduction. This metric quantifies reduction in P4 installation state only; it does not quantify traffic reduction, Snort rule reduction, execution-time reduction, or throughput improvement.

Provenance-based candidate rule sets.

Backend lowering retains the `source_rule_ids` provenance of each final P4 entry. Let $E(x)$ be the set of retained P4 entries whose predicates match a parser-eligible packet x, and let $S(e)$ be the source Snort rule IDs associated with entry e. The candidate rule set is

$$C(x) = \bigcup_{e \in E(x)} S(e).$$

Thus, multiple source rules may map to one final entry, a packet may match multiple entries, and duplicate IDs are removed by set union. For the strong-discriminator ablation of Section 7.6, the P4-eligible baseline is $B = 739$ rules, and per-packet candidate-rule-space reduction is

$$R(x) = \left(1 - \frac{|C(x)|}{B}\right) \times 100\%.$$

This metric is defined only within the 739-rule P4-eligible partition. The 3,257 operational fallback rules and 21 semantic UNSUPPORTED fallback rules are outside B. A parser-eligible packet with no retained predicate match can have an empty $C(x)$; this is not interpreted as safe full packet bypass because fallback rules remain applicable. Because those 3,278 rules apply to every packet, we also report the reduction over the full 4,017-rule ruleset,

$$R_{\text{full}}(x) = \left(1 - \frac{|C(x)| + 3278}{4017}\right) \times 100\%,$$

which can never exceed 18.40%. Under the earlier compilation the same definitions used $B = 3611$ and 406 fallback rules.

For a grouping of the retained entries, a prefilter that forwards whole groups passes Snort every rule of each group containing a matched entry. Its grouped candidate set $C_G(x)$ is the union of the rules of those groups, and the same two reductions are computed with $C_G(x)$ in place of $C(x)$.

Observed preservation measures. We distinguish four finite, empirical checks. *Semantic split preservation* compares FULL Snort alerts with the union from semantic candidate and UNSUPPORTED fallback rules, reporting missing and additional alerts. *P4 alert-packet preservation* checks whether observed parser-eligible packets associated with semantic-candidate alerts are marked P4_CANDIDATE, reporting candidate alert packets, parser eligibility, P4 candidates, and P4 no-matches. *Strong-discriminator observed-event coverage* checks retained-predicate coverage for P4-eligible alert events while accounting for operational and semantic fallback events. Finally, *candidate-set SID membership* checks whether an observed P4-eligible alert SID belongs to $C(x)$. The last measure does not validate dynamic stateful Snort execution using $C(x)$, because prior packets, stream state, flowbits, and other stateful semantics may be required.

6.4 Validation Procedure and Held-Out Protocol

The evaluation includes a controlled validation and CICIDS2017 [13] Thursday web-attack windows. SQLi and XSS are development/diagnostic windows; Brute Force is reserved as held-out evaluation. Frozen policy and grouping configurations are evaluated on Brute Force without retuning after its inspection. The initial scenarios considered controlled cases designed to exercise the semantic

classes produced by the compiler. Rather than attempting to evaluate the complete Snort 3 rule language, the experiments focus on representative rules that expose different relationships between the original IDS rule and the generated data-plane predicate.

The first scenario is an early semantic-model unit scenario for a header-only EXACT-class expectation. This is a component-level test specification, not a trace-based EXACT result for the frozen Community ruleset. That ruleset produced EXACT= 0; the reported experiments therefore do not claim empirical EXACT validation.

The second scenario considers rules that combine supported header predicates with unsupported conditions, such as payload content. These rules are expected to produce NECESSARY representations. The generated data-plane predicates are used only for pre-filtering, while Snort retains responsibility for evaluating the remaining conditions and producing the final detection decision.

The third scenario considers rules for which the initial target profile cannot derive a useful and safe pre-filtering predicate. Such rules are classified as UNSUPPORTED and remain entirely delegated to Snort. This scenario verifies that unsupported semantics are explicitly reported rather than silently removed during compilation.

6.5 End-to-End Differential Campaign

The trace-based checks exercise only the few rules that fire in the evaluated windows. To test the safety relation on a wider set of rules, a differential campaign uses Snort 3 as the oracle and BMv2 as the system under test. For each rule, a generator synthesizes from the rule's IR a short flow whose witness packet should make Snort raise that rule's alert: a TCP handshake when the `flow` option requires an established session, then a witness that carries the rule's header values, places each `content` in the inspected buffer, and uses a service port that lets Snort bind the required inspector. Each witness has two neighbors: a payload-negative frame with the same 5-tuple but without the content, on which Snort is expected not to alert, and, where the installed entries leave one, a header-negative frame whose destination port falls outside every installed entry.

A witness could be synthesized for 2,142 of the 4,017 rules. The other 1,875, including the 14 rules that both compilations classify as UNSUPPORTED, were excluded: 1,654 of them carry options such as `pcre`, `byte_test`, `byte_jump`, or `flowbits` that cannot be satisfied by choosing bytes, and the rest have content the generator cannot place or a header it cannot resolve. A sample of 1,000 rules, stratified proportionally by protocol, service, and direction with a fixed seed, was written to a trace in which each case occupies its own second: 6,684 frames, all inside D. The same sample was regenerated whenever the witness generator was revised.

Snort 3 reads the trace with the full ruleset and the experimental profile, and each alert is attributed to a case twice, by packet range and by timestamp. BMv2 receives the same frames one at a time, and the counter category that increments after each frame is recorded. Each case then receives one verdict.

It is CONFIRMED when Snort alerted and both the witness and every frame it alerted on were marked candidate. It is a VIOLATION when Snort alerted but the witness, or another frame that satisfies the rule's header and on which Snort alerted, was marked no-match; this is the only outcome that contradicts $r \Rightarrow$ cand. It is VACUOUS when Snort did not alert, so the case did not test the implication. It shows STATE DEPENDENCE when the witness was a candidate but Snort attributed the alert to a session frame that does not satisfy the rule's header and that the data plane marked no-match. Rules placed in operational fallback by the strong-discriminator policy are reported separately, because the prefilter is not responsible for them.

The same trace was run against the 235-entry global plan and the 225-entry strong-discriminator plan. In the global plan the 21 UNSUPPORTED rules form the fallback; in the strong-discriminator plan, the 3,278 rules outside B. The global plan limits what the campaign can falsify: entry `p4-88` matches every in-domain TCP packet, so a TCP witness is a candidate however its own rule was lowered. To quantify this, the offline matcher also records whether each witness is matched only by entries whose provenance contains the case's own rule. For such a case, a lowering defect in that rule would have produced a VIOLATION.

The campaign places each witness on a port of the rule's own header, so it cannot exercise the service-based evaluation of Section 4.5. A second, dedicated campaign does: it takes 200 HTTP rules with `service:http`, sampled from the synthesizable ones, and moves the HTTP flow to destination port 18080, which none of their headers lists (Section 7.2).

7 Results

7.1 Semantic Compilation Coverage (RQ1)

The compilation processed and successfully parsed all 4,017 rules in the evaluated Snort 3 Community ruleset. The semantic manifest contains no EXACT classifications, 3,996 NECESSARY classifications, and 21 UNSUPPORTED classifications. Hence, 99.48% of the rules contributed a candidate predicate under the supported profile. For 1,072 of them, the 1,060 port-based rules and the 12 that name a service protocol in their header and keep their port checks, that predicate is derived from the full header. For the other 2,924, which are service-based, it keeps only the header addresses, and the TCP protocol where a `flags` option requires it (Section 4.5).

This percentage characterizes semantic candidate coverage, not complete offloading of Snort semantics. For a rule r and its NECESSARY predicate p_r, the relation the compiler targets is $r \Rightarrow p_r$; neither $p_r \Rightarrow r$ nor $r \Leftrightarrow p_r$ is assumed. Payload inspection, stream and flow/state semantics, normalization, and other target-inexpressible conditions remain delegated to Snort. For the 21 UNSUPPORTED rules, no safe selective P4 predicate was established, and the rules remain in the Snort fallback rather than being ignored. The absence of

EXACT rules therefore reflects a deliberate conservative classification policy: the compiler does not declare equivalence when it has preserved only a necessary condition.

The manifest explains the 21 UNSUPPORTED cases without silently discarding them. Four are traditional rules whose represented headers are nonselective while payload or byte-level conditions remain delegated; ten are service-style rules whose service, flow, buffer, and payload semantics are delegated and for which the parser intentionally invents no traditional packet headers. The other seven are service-based rules with unrestricted addresses, SIDs 1941 (TFTP), 28852, 29174, and 38993 (HTTP), and 56583 to 56585 (Kerberos), whose only selective conditions were the ports and protocol that Snort does not check once it has identified their service. Thus, each lacks a selective represented H_r and remains in the Snort fallback. Across NECESSARY rules, the manifest commonly records content and other payload/buffer conditions, flow or state conditions, application/service semantics, and PCRE as delegated while retaining at least one selective represented header predicate. A smaller set also records a fallback header form, such as a port range, negated port, or unsupported flag encoding, but is emitted only when another selective represented predicate is available.

Answer to RQ1. Under the evaluated target profile and within the enforced packet domain, the compiler derives explicit necessary predicates for 3,996 of the 4,017 rules while making the remaining semantic limitations visible; traffic outside that domain is retained by failing open rather than by an unstated parsing assumption. For the 2,924 service-based rules among them, a predicate that stays necessary once Snort identifies the flow's service keeps only the header addresses, which costs selectivity rather than coverage (Section 7.2).

7.2 Service-Based Rules and the Earlier Compilation

Of the 4,017 rules, 2,928 declare a `service` option, most often `http` (2,314), `imap` (150), `pop3` (138), and `smtp` (131). Twelve more name `http` in their header, which gives them that service but keeps their port checks. Three, SIDs 529, 3238, and 3239, are service-based without a `service` option: their `dce_iface` and `dce_stub_data` options add `dcerpc` and `netbios-ssn` (Section 4.5). Service-based rules number 2,931 in all; 1,060 rules use no option that gives them a service.

An earlier version of the compiler represented every rule by its full header and recorded a service as one more delegated condition. That would be sound if Snort required the service and the header together; Snort 3 instead replaces the port and protocol checks by the service match. Every predicate of the 2,928 rules with a `service` option constrained the protocol, and 2,877 also constrained a port, so a rule's own predicate could be false on a packet it alerts on. Table 4 compares that compilation with the current one; its artifacts are archived and reproduce byte for byte.

Whether an over-constrained predicate becomes a missed candidate depends on the whole plan. In the earlier 331-entry global plan, two port- and protocol-free entries, which match a destination or a source in `HOME_NET`, contained the

Table 4: Earlier and current compilation of service-based rule headers. Trace values are for held-out Brute Force; campaign and off-port values are BMv2 measurements under the strong-discriminator plan of each compilation.

	Earlier (full header)	Current (addresses only)
Necessary / Unsupported / Exact	4,003 / 14 / 0	3,996 / 21 / 0
Predicate instances / final entries	6,675 / 331	4,040 / 235
Strong plan: eligible / fallback rules	3,611 / 406	739 / 3,278
Strong plan: entries	322	225
Strong plan: candidate rate	98.942%	94.885%
Mean candidate set (rules)	324.52	4.48
Full-ruleset reduction	81.81%	18.29%
Campaign: confirmed / fallback / state dep.	851 / 107 / 2	134 / 826 / 0
Off-port HTTP rules: violation / fallback	190 / 0	0 / 190

region on which Snort can alert once the service is identified for 2,922 of the 2,931 rules. A TCP-only entry contained the TCP part of that region for the other nine: the seven rules above that are now UNSUPPORTED, and SIDs 1415 and 1416 (SNMP over UDP to the broadcast address). That plan was therefore candidate-complete for these rules except for non-TCP packets of those nine. Its strong-discriminator ablation removed the broad entries, and none of the 322 entries it kept contained the region of any of the 2,931 rules, 2,877 of which were among its 3,611 P4-eligible rules.

The dedicated off-port test of Section 6.5 exercises this case. Of its 200 HTTP rules, Snort alerted for 190, 190 distinct SIDs, on port 18080, so it did not check the header port; the other 10 raised no alert. With the earlier strong-discriminator plan installed on BMv2, all 190 witnesses were no-match, and so was the flush acknowledgment to which Snort attributes each alert: a prefilter that withheld no-match packets would have dropped both. The earlier global plan marked all 190 witnesses candidates, through two port-free entries of other rules. Under the current compilation there is no such case. The global plan marked all 190 witnesses candidates, each matched by an entry of its own rule as well as by others; the strong-discriminator plan places all 190 rules in operational fallback, which Snort inspects without the prefilter, so the 60 witnesses it marked no-match violate nothing. The test covers one port and one service; the generalization rests on the source reading of Section 4.5.

The correction moves selectivity rather than coverage. Without ports and protocol, most service-based rules keep only a `HOME_NET` address condition, which the strong-discriminator policy does not treat as discriminating (Section 7.6), so the policy keeps 739 rules P4-eligible instead of 3,611. Much of the candidate-rule-space reduction that the earlier compilation appeared to offer came from header ports that Snort 3 does not check once it knows the service (Section 7.7).

7.3 P4 Backend Feasibility (RQ2)

Table 5 summarizes the semantic compilation and backend feasibility results. The backend generated, expanded, and normalized 4,040 predicate instances. These instances yielded 231 unique matches before TCP-flag lowering. Two flag-constrained groups required expansion, adding four entries and resulting in 235 final installation entries. The P4Runtime plan encoded 235/235 entries, with zero encoding failures and zero validation errors, within the 8,192-entry table capacity. On BMv2, every installation of the 235-entry plan, and of the 225-entry strong-discriminator plan, read back all of its entries.

The reduction from 4,040 source instances to 235 final entries (94.18%) is an *installation-state reduction*: identical backend matches are merged before installation. Most of it comes from the service-based rules, whose address-only predicates coincide: a single entry, matching a destination in `HOME_NET`, carries the provenance of 2,043 rules. The reduction does not represent traffic reduction, elimination of Snort rules, execution-time improvement, or speedup.

Answer to RQ2. The semantically derived predicates were lowered to the evaluated target within its table capacity and P4Runtime validation constraints, without concealing unsupported semantics.

Table 5: Semantic compilation and backend feasibility. The 94.18% value is backend installation-state reduction, not traffic reduction.

Rules	Exact	Necessary	Unsupported	Predicate instances	Unique	Final	P4Runtime
4017	0	3996	21	4040	231	235	235/235

7.4 Packet-Domain Enforcement and Validation

Because the selectivity and candidate-set metrics are defined over the enforced domain D, the correspondence between the deployed parser and the offline matcher was validated before those metrics are reported.

The controlled data-plane suite passed all eleven cases with exact counter deltas. TCP, UDP, and ICMP candidates, a guaranteed no-match, and exact and masked TCP flags were classified inside D as expected. IPv4 options, a first fragment with MF set, a non-initial fragment, a truncated IPv4 header, and a truncated TCP header were each routed to `PACKET_DOMAIN_FALLBACK`, with no candidate or no-match count. No case produced a data-plane verdict that disagreed with the offline expectation.

Across the evaluated traces, 350 of 35,274 SQLi frames, 2,772 of 191,848 XSS frames, and 7,199 of 1,108,658 Brute Force frames fell outside D. These are dominated by non-IPv4 EtherTypes: ARP (125, 1,660, and 4,074), IPv6 (195, 820, and 2,523), and smaller LLDP and non-IP EtherType populations. Only the held-out trace contains direct-IPv4 frames excluded by the header gate itself: 16 frames with fragmentation or reserved flags. They are the reason its

Table 6: Packet-level selectivity, measured on BMv2 by replaying each trace whole under each plan, with exact accounting; the offline classifier gives the same counts. Candidate rates are among parser-eligible packets, that is, packets inside the enforced domain D; frames outside D are retained as candidates through domain fallback and are excluded from the denominator.

Trace	Total	Eligible	Global candidate	Global no-match	Strong candidate	Strong no-match
SQLi	35274	34924	34924	0	92.298%	7.702%
XSS	191848	189076	189076	0	92.173%	7.827%
Brute Force (held-out)	1108658	1101459	1101459	0	94.885%	5.115%

parser-eligible population is 1,101,459 rather than the 1,101,475 obtained when fragmentation is not gated; all 16 remain candidates through domain fallback.

Trace replay accounting was reconciled against the data plane. Each of the three traces was replayed whole under the global and the strong-discriminator plans, and in all six replays the switch counted every frame sent, with each counter category equal to the offline classifier's value (Table 6). The replay protocol's longer veth queues and 2 ms inter-frame delay come from an earlier held-out replay without them, which lost 87 of 1,108,658 frames (0.008%) in the transport layer, all within the candidate category.

7.5 Global Binary Prefilter Selectivity (RQ3)

Table 6 reports packet-level selectivity under the current Ethernet/IPv4 parser scope and supported protocols, measured on BMv2 by replaying each trace whole. The global candidate bit marked 100% of parser-eligible packets as candidates in SQLi, XSS, and the held-out Brute Force trace. Parser-ineligible frames are outside this denominator and are not treated as P4 no-matches. Thus, within the evaluated CICIDS2017 windows, high semantic candidate coverage did not translate into useful packet-level bypass selectivity for the global binary candidate bit.

7.6 Selectivity Causal Analysis

The causal audit first identified broad NECESSARY predicates in the global table. Entry `p4-88` matches TCP protocol only; it subsumes 162 of the 235 installation entries (68.94%). Entry `p4-10` matches traffic with destination `192.168.10.0/24` and subsumes 154 entries (65.53%), whereas `p4-51` matches the corresponding source prefix and subsumes 71 entries (30.21%). The last two carry the provenance of 2,043 and 900 rules, almost all of them service-based rules whose predicates keep only a `HOME_NET` address. Together, these three entries cover every parser-eligible packet of SQLi and XSS; `p4-88` alone covers 84.87% and 88.54%. Deduplication is not responsible for the coverage, because it merges identical normalized matches rather than removing restrictions; TCP-flag

lowering is also not responsible, because these broad entries do not depend on TCP flags.

The strong-discriminator ablation treats an installation entry as a strong discriminator when it constrains a source or destination port, the TCP flags, or an IPv4 prefix other than exactly `HOME_NET`; a rule remains P4-eligible only if every final entry derived from it is strong. The criterion is structural and uses no traffic frequencies. It removed ten entries, each constraining at most the protocol and exactly `HOME_NET` as source, destination, or both, `p4-10`, `p4-51`, and `p4-88` among them, retaining 225 installation entries for 739 P4-eligible rules. It places 3,257 semantic candidates in operational fallback and retains the 21 semantic UNSUPPORTED rules in fallback; it is therefore an operational ablation, not a revised semantic classification. Its plan was installed on BMv2 for the trace replays and for the end-to-end campaign of Section 7.10. As shown in Table 6, the candidate rate nevertheless remained between 92.173% and 92.298% on the development traces and 94.885% on the held-out trace, where 56,343 parser-eligible packets were no-matches.

The residual-selectivity audit shows what limits the retained union. Four retained entries reach 90% of the parser-eligible packets of SQLi and XSS. The first matches TCP to `HOME_NET` with the ACK bit set and covers 46.80% and 48.87% by itself; it comes from a single service-based SMTP rule, SID 309, whose `flags:A+` option keeps its TCP protocol and adds the ACK bit. The other three match traffic from `HOME_NET` to destination ports 443, 80, and 53. All four are strong discriminators by the structural criterion, but their flag and port conditions coincide with the dominant traffic of the workload. Under the earlier compilation, the same ablation still marked 98.869–99.015% of the packets, and 22 retained entries were needed to reach 99% of SQLi: a distributed union rather than a few dominant entries. Either way, semantic safety and operational selectivity are distinct properties: a predicate can be sound as a necessary condition, and structurally specific, while adding little packet-level selectivity when composed with the rest of the ruleset.

Answer to RQ3. The global binary prefilter is conservative but operationally non-selective on the evaluated traces. The strong-discriminator ablation, which removes the broad entries, still marks 92.2–94.9% of parser-eligible packets as candidates.

7.7 Candidate Rule-Set Reduction (RQ4)

The global candidate bit asks whether a packet remains a candidate for IDS handling. The provenance-based analysis asks a different question: which of the 739 P4-eligible Snort rules remain plausible for a given packet? For each parser-eligible packet, it unions the source-rule IDs of the matched retained P4 entries to derive a per-packet candidate rule set.

The mean candidate set contained 3.97 rules for SQLi, 3.94 for XSS, and 4.48 for held-out Brute Force, a mean reduction of 99.39–99.47% within the P4-eligible partition (Table 7), and 99% of packets had at most eight candidate rules. Brute Force was not used for policy retuning. The partition is small,

Table 7: Offline per-packet candidate rule set feasibility using the strong-discriminator baseline of 739 P4-eligible rules. The last column counts the 3,278 fallback rules, which apply to every packet, against the full 4,017-rule ruleset; it cannot exceed 18.40%.

Trace	Base.	Mean	Median	P90	P95	P99	Max	Mean reduction	Full-ruleset reduction
SQLi	739	3.97	3	8	8	8	43	99.46%	18.30%
XSS	739	3.94	3	8	8	8	43	99.47%	18.30%
Brute Force (held-out)	739	4.48	3	8	8	8	293	99.39%	18.29%

however. The 3,278 fallback rules apply to every packet, and the mean reduction over the full ruleset is 18.30% for SQLi and XSS and 18.29% for held-out Brute Force, just under the 18.40% that no candidate set can exceed.

The earlier compilation gave a very different picture (Table 4). Its strong-discriminator ablation kept 3,611 rules P4-eligible, the held-out mean candidate set held 324.52 of them, and the reduction over the full ruleset was 81.81%. Most of the difference is the header ports of service-based rules: with them, a packet to port 80 excluded the rules written for other ports; without them, those rules are in fallback and apply to every packet. Header provenance cannot recover that discrimination for rules whose header ports Snort 3 does not check. The identified service, which Snort itself uses to reject rules of other services, could.

This is an offline candidate-rule-space feasibility measurement, not a measured reduction in Snort execution, throughput, or latency, and not a count of rules dynamically executed by Snort. A candidate-set-aware Snort execution engine was not implemented.

7.8 Coarse-Grained Grouping Ablation (RQ4)

A prefilter could also forward groups of rules rather than per-entry provenance. We assigned the 225 retained entries to groups by protocol (4 groups), protocol plus direction (12), protocol plus service (203), and protocol plus direction plus service (213), and measured the grouped candidate set on the three traces (Table 8). For each parser-eligible packet it is the union of the rules of every group that contains a matched strong-discriminator entry, which is what a prefilter that forwards whole groups would pass to Snort.

Under protocol grouping, the TCP group holds 642 of the 739 rules, and the median packet activates it. The grouped set averages 551 to 590 rules, a reduction of 20.10–25.38% within the P4-eligible partition and of only 3.70–4.67% over the full ruleset. Adding direction reaches 57.00–60.11%. Protocol plus service, whose largest group holds the 291 rules of destination port 1521, reaches 98.80–98.96%, and protocol plus direction plus service 98.98–99.08%, within half a percentage point of the entry-level reduction; over the full ruleset, both stay within 0.12 points of it. Under the earlier compilation, with a P4-eligible partition five times larger, the cost of coarse grouping was higher: the TCP group held 3,439 of 3,611 rules, and protocol grouping reduced the set by only

Table 8: Grouped candidate set on the traces: mean rules forwarded per parser-eligible packet when every rule of each activated group is forwarded, and the reduction over the 739 P4-eligible rules (strong-discriminator plan, offline). The first row is the entry-level set of Table 7.

		SQLi		XSS		Brute Force	
Strategy	Groups	Mean	Red.	Mean	Red.	Mean	Red.
Entry level	–	3.97	99.46%	3.94	99.47%	4.48	99.39%
Protocol	4	551.44	25.38%	562.04	23.95%	590.47	20.10%
Protocol, direction	12	294.77	60.11%	307.89	58.34%	317.77	57.00%
Protocol, service	203	7.71	98.96%	7.73	98.95%	8.89	98.80%
Protocol, direction, service	213	6.78	99.08%	6.89	99.07%	7.52	98.98%

12.30–19.10%, against 91.01–96.11% at entry level.

The reductions in Table 7 therefore come from the mapping of matched P4 entries to their source-rule IDs, followed by the union of those provenance sets for each packet, and grouping preserves them only when its groups are nearly as fine as the entries.

7.9 Dynamic Alert Preservation (RQ4)

Table 9 separates three of the empirical properties defined in Section 6: semantic split preservation, P4 alert-packet preservation under each plan, and observed-event coverage under the strong discriminator. The fourth, candidate-set SID membership, has no P4-eligible alert event to test under this compilation. These are not interchangeable claims of equivalence.

In SQLi, XSS, and held-out Brute Force, FULL Snort produced 150, 528, and 2,391 alerts from one, two, and five SIDs: 254; 254 and 402; and 254, 399, 402, 650, and 1394. The semantic split had zero missing or additional alerts. It was measured by running Snort separately on the candidate and UNSUPPORTED partitions of the earlier compilation. The seven rules that the current compilation adds to UNSUPPORTED raised no alert in the FULL runs and take part in no flowbit dependency, so the result carries over unchanged. No UNSUPPORTED rule fired in any of the three traces, which limits the discriminating power of this comparison; Section 9 states this explicitly.

P4 alert-packet preservation was measured on BMv2 by replaying only the alert packets of each trace under each plan, sending the captured bytes unchanged, with exact accounting. Under the global plan, all 150, 528, and 2,391 alert packets were candidates. Under the strong-discriminator plan, all but three were no-match. This is not a loss, because every observed alert comes from a rule in operational fallback, which Snort inspects without the prefilter. SID 254, which raised 2,960 of the 3,069 alerts, declares the `dns` service; without its UDP and port-53 conditions its predicate keeps only the `HOME_NET` destination, which is not a strong discriminator. The offline strong-discriminator audit accordingly attributes all 3,069 alert events to fallback and none to retained predicates, so the traces do not exercise the candidate sets of Table 7 at all. Under the earlier

compilation, 2,289 of the held-out events were covered by retained predicates, and in all 2,289 the alert SID belonged to the packet's candidate set.

Thus, observed alert preservation was 100% in the three traces. The trace names do not imply that the observed Snort alerts detected the corresponding CICIDS2017 attack labels. These measurements describe the evaluated ruleset behavior and constitute observed alert preservation, not formal equivalence, guaranteed preservation for arbitrary traffic, or detection accuracy.

Table 9: Observed alert preservation. Missing and additional compare the semantic split with FULL Snort. P4 columns concern the alert packets, replayed alone on BMv2 under each plan with exact accounting; under the strong-discriminator plan every alert belongs to a fallback rule.

Trace	FULL	Missing	Add.	Alert packets	Global cand.	Strong cand.	Strong no-match	Via fallback	Pres.
SQLi	150	0	0	150	150	0	150	150	100%
XSS	528	0	0	528	528	0	528	528	100%
Brute Force (held-out)	2391	0	0	2391	2391	3	2388	2391	100%

7.10 End-to-End Differential Campaign (RQ4)

Table 10 reports the campaign verdicts. No case produced a VIOLATION under either plan, and no frame fell outside D, so no case was resolved by domain fallback. The two attributions of Snort alerts to cases, by packet range and by timestamp, agreed for every case. In both runs, the category predicted by the offline matcher agreed with the BMv2 counter on all 2,101 witness and neighbor frames, which extends the eleven-case packet-domain suite of Section 7.4 to a thousand rules and to the matcher that produced the selectivity results.

Under the global plan, 958 of the 1,000 rules were CONFIRMED, 40 were VACUOUS, and 2 were in fallback: SIDs 56583 and 56584, Kerberos rules that are UNSUPPORTED in this compilation. All 1,000 witnesses and all 1,000 payload-negative frames were candidates, and all 101 header-negative frames were no-matches. Every confirmed witness was matched by its own rule's entries but also by an entry of another rule, so no confirmation under this plan can be attributed solely to the rule's own predicate. This is the broad coverage of Section 7.6 seen at the level of individual rules: the same coverage that removes packet-level selectivity also removes the power to falsify a single rule's lowering.

Under the strong-discriminator plan, 134 rules were CONFIRMED, 826 were in operational fallback, 40 were VACUOUS, and none showed STATE DEPENDENCE. Every confirmed witness was matched by at least one entry of its own rule, and for 36 of the 134 it was matched by no other entry, so a lowering defect in those rules would have been observed as a VIOLATION. Of the 826 fallback rules, 78 had witnesses marked no-match, as expected for rules with no retained entry, and 748 were candidates through other rules' entries. The payload-negative frames measure the conservative excess of a header-level prefilter: 915 were candidates on which Snort did not alert, and 7 were candidates on which Snort did alert,

because their rules carry no `content` to remove. The remaining 78 belong to fallback rules, which the plan does not prefilter.

The campaign also shows how much HTTP detection depends on packets that do not satisfy a rule's header. For 606 of the fallback rules, all HTTP rules, Snort attributed the alert to the acknowledgment from the other endpoint that triggers inspection of the buffered HTTP message, and the strong-discriminator plan marked 445 of those acknowledgments no-match. Nothing is lost, because these rules are in fallback. Under the earlier compilation, the same mechanism produced two STATE DEPENDENCE cases, SIDs 43809 and 47386, HTTP request rules that its strong-discriminator plan prefiltered. The compiled predicate held on the witness, but a deployment that withheld no-match packets from Snort could have delayed or suppressed those alerts. The global plan hides this dependence, because `p4-51` matches every packet leaving `HOME_NET`.

Table 10: End-to-end differential campaign over one stratified sample of 1,000 rules, replayed as the same trace under both plans. *Fallback* counts rules the plan does not prefilter: the UNSUPPORTED rules, and under the strong-discriminator plan also its operational fallback. *Own-only* counts confirmed rules whose witness was matched only by entries of their own rule, computed with the offline matcher.

Plan	Entries	Confirmed	Violation	State dep.	Fallback	Vacuous	Own-only
Global	235	958	0	0	2	40	0
Strong discriminator	225	134	0	0	826	40	36

Answer to RQ4. The evaluated predicates preserved all observed alerts in the three traces, as candidates under the global plan and through fallback under the strong-discriminator plan. In the end-to-end campaign no case was a VIOLATION: no rule that a plan prefilters had its witness marked no-match while Snort alerted, across 958 confirmed rules under the global plan and 134 under the strong-discriminator plan, 36 of the latter attributable to the rule's own entries alone. The predicates also exposed per-packet candidate rule sets of about four rules within the 739-rule P4-eligible partition, but the reduction over the full ruleset is 18.29–18.30%, because the other 3,278 rules apply to every packet. That result remains an offline feasibility measurement, because dynamic stateful candidate-set execution was not implemented, and the traces exercise no retained predicate. A prefilter that forwards whole groups keeps the reduction only with groups nearly as fine as the entries.

Figure 2 summarizes the semantic compilation flow and distinguishes the implemented binary prefilter from the offline provenance analysis.

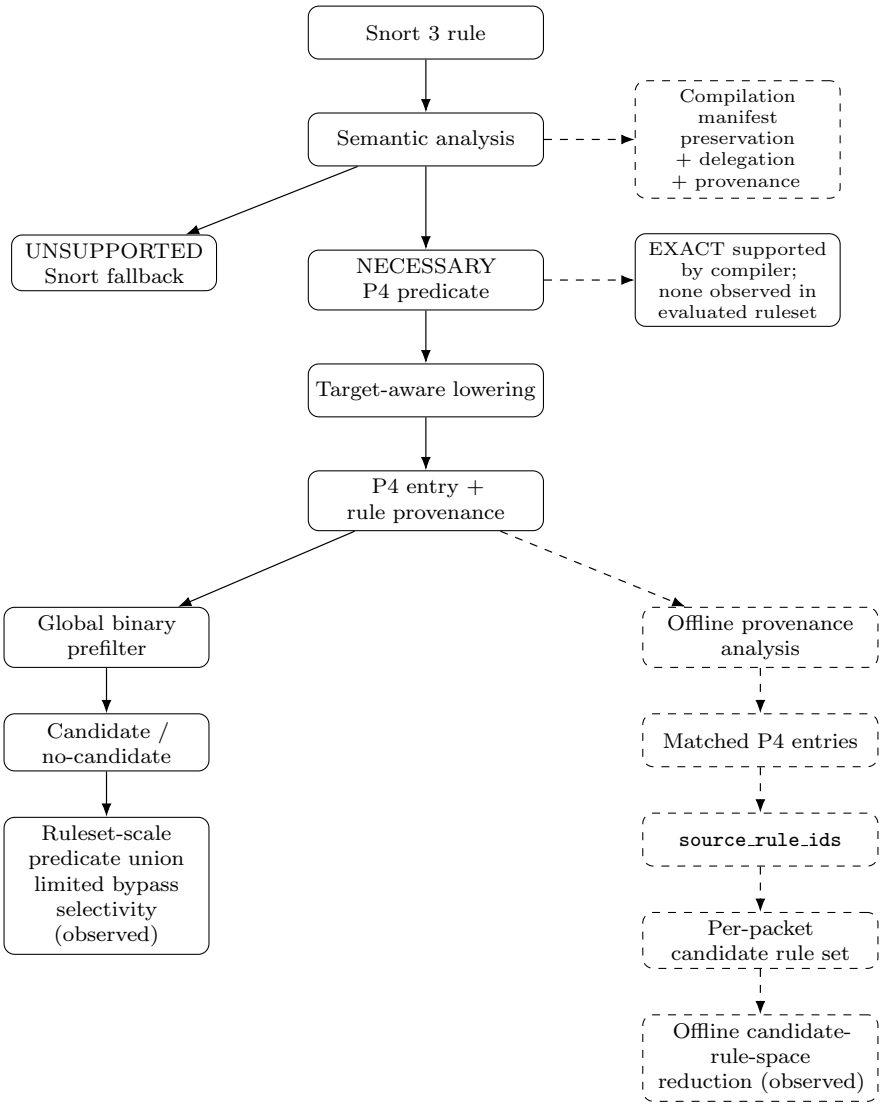


Figure 2: Semantics-aware, target-aware compilation overview. Semantic analysis retains unsupported or delegated semantics in Snort while deriving necessary P4 predicates; for service-based rules they keep only the header addresses (Section 4.5). Target-aware lowering preserves source-rule provenance. The implemented global binary prefilter aggregates predicates into candidate metadata and showed limited packet-bypass selectivity in the evaluated ruleset. The dashed provenance branch is an offline feasibility analysis for deriving per-packet candidate rule sets; dynamic candidate-set-aware IDS execution is not implemented. EXACT is supported by the classifier but was not observed in the evaluated ruleset.

8 Discussion

8.1 Semantic Safety Does Not Imply Operational Selectivity

For a Snort rule r and a generated P4 predicate p_r, the relation $r \Rightarrow p_r$ states that p_r is a safe necessary condition under the supported semantic profile and inside the enforced packet domain D, with traffic outside D retained through fail-open fallback. It does not establish $p_r \Rightarrow r$ or $r \Leftrightarrow p_r$. Semantic safety therefore asks whether traffic that satisfies the rule also satisfies its compiled predicate; operational selectivity asks how much other traffic that predicate can exclude.

Under the evaluated profile, 3,996 of 4,017 rules were classified as NECESSARY, 21 as UNSUPPORTED, and none as EXACT. Nevertheless, the global binary prefilter marked 100% of parser-eligible packets as candidates in SQLi, XSS, and held-out Brute Force. This is not a contradiction. The usefulness of a safely compiled predicate is not determined solely by whether it preserves a necessary condition, but also by how that predicate composes with predicates generated for the remainder of the ruleset. The absence of EXACT classifications is both a limitation of the evaluated target/profile and evidence that partial semantic preservation was not promoted to equivalence.

8.2 The Union Effect in Large Rule Sets

Conceptually, the global candidate decision is the disjunction $P_{\text{global}}(x) = p_1(x) \vee p_2(x) \vee \cdots \vee p_n(x)$. Individual predicates may be useful filters, but the aggregate coverage of a large disjunction can approach the workload population. The causal audit identified broad initial contributors: `p4-88` matches TCP only, `p4-10` matches a destination in `HOME_NET`, and `p4-51` matches a source in `HOME_NET`; the last two carry the provenance of most service-based rules.

Broad predicates were an initial contributor, but they were not a complete explanation. The strong-discriminator ablation retained 225 entries for 739 P4-eligible rules, placed 3,257 rules in operational fallback, and retained 21 semantic UNSUPPORTED rules in fallback. Even then, the held-out Brute Force trace had a 94.885% candidate rate among parser-eligible packets. A few retained entries are structurally specific, an ACK bit toward `HOME_NET` and the web and DNS ports from it, yet coincide with most of the traffic; under the earlier compilation, the residual coverage was instead spread across many retained predicates. Thus, the evidence does not attribute poor selectivity to only three broad entries. It instead indicates that large heterogeneous rulesets, together with the traffic they are written for, can make a single global candidate bit too coarse and discard rule-level provenance.

The union also limits what can be tested. In the end-to-end campaign, every confirmed witness under the global plan was also matched by another rule's entry, so a defect in any single rule's lowering would have been masked by the rest of the table. Only after the broad entries were removed did 36 confirmations

depend on the rule's own entries alone. Testing the lowering of an individual rule therefore needs a configuration in which that rule's predicate is not covered by the others, a condition closely related to the one that packet-level selectivity depends on.

8.3 From Packet Bypass to Candidate-Rule Selection

The global prefilter asks whether a packet should remain a candidate for IDS handling. Provenance enables a different question: which IDS rules remain plausible candidates for that packet? The latter is obtained by mapping matched P4 entries to their recorded source-rule IDs and taking their union as a per-packet candidate rule set. Rule-level provenance consequently retains information that a global candidate bit discards.

Using the 739-rule strong-discriminator baseline, the mean candidate set contained about four rules in each trace, 4.48 on held-out Brute Force (median 3; P99 8; maximum 293), a reduction above 99.3% within that partition. The partition, however, holds only the rules whose predicates the policy treats as discriminating. The 3,278 fallback rules bound what any candidate set can save: counted against the full ruleset, the reductions are 18.29–18.30%, and no packet can fall below 3,278 applicable rules.

The comparison with the earlier compilation locates the loss. With the header ports of service-based rules, the same derivation reduced the held-out candidate set by 81.81% over the full ruleset; without them, most of those rules keep only an address condition on `HOME_NET` and fall back. Much of what header provenance appeared to discriminate was therefore the header port, which Snort 3 does not use once it knows the service. A candidate selector that also used the identified service, which Snort itself uses to reject rules of other services, is a natural next step; this paper does not evaluate one.

This observation motivates a future design space involving rule masks, rule partitions, hierarchical selectors, candidate metadata, and dependency-aware scheduling in hybrid P4/IDS systems. It is not evidence of a runtime implementation: candidate-set-aware Snort execution was not implemented, no throughput or latency benchmark was conducted, and no speedup is demonstrated. In particular, dynamic stateful rule execution remains unresolved.

8.4 Why Coarse Grouping Does Not Help

The candidate-rule-space reduction within the P4-eligible partition was produced by fine-grained entry-to-source-rule provenance, not by grouping. Measured directly, grouping gives back the discrimination it aggregates: a coarse group forwards every rule it contains, and under protocol grouping, whose TCP group holds 642 of the 739 rules, the mean reduction falls to 20–25%. Finer groupings recover it step by step, and the two that include the service come within half a percentage point of entry-level provenance (Table 8). Over the full ruleset, the fallback rules dominate either way.

Coarse structural groups aggregate rules whose necessary predicates remain different. This does not show that every grouping is unhelpful: the four evaluated strategies are structural, and the grouped set was computed offline. Future grouping designs could account for provenance, semantic dependencies, rule overlap, flow/state dependencies, and hardware resource constraints.

8.5 Implications for P4-Based IDS Offloading

The evaluated results suggest that translation rate alone is insufficient for assessing a P4-based IDS prefilter. Classifying 99.48% of rules as NECESSARY did not imply operationally selective binary packet bypass. This complements prior work on rule-derived P4 filtering and resource-aware data-plane representations [7, 14]: the semantic relation between the original rule and the compiled predicate must remain explicit.

Semantic transparency should therefore be a first-class compiler output. A useful artifact records preserved predicates, delegated predicates, fallback rules, and unsupported rules rather than presenting partial translation as full rule offload. The evidence also supports evaluation along multiple axes: semantic coverage, target feasibility, installation state, packet-level selectivity, observed alert preservation, and candidate-rule-space reduction. Work on context required by rule-based prefiltering further reinforces the need to account for IDS state beyond packet-header matching [3].

The service-based case shows that even the header of a rule is not self-describing. What a header condition means depends on how the IDS evaluates it, and Snort 3 drops the port and protocol conditions of service-based rules once a flow's service is known. The earlier compilation, which derived port conditions from those headers, produced a strong-discriminator plan that marked no-match packets on which Snort alerted; its global plan stayed candidate-complete for all but nine of these rules only through the broad entries that removed packet-level selectivity. Compiled as Snort 3 evaluates them, the rules are safe under both plans, but their predicates shrink to address conditions on `HOME_NET` that match most of the traffic, and the strong-discriminator policy leaves them to fallback. Safety and selectivity pulled against each other within the same rules.

Binary packet bypass may be too coarse for large heterogeneous rulesets under the evaluated configuration. Predicate provenance can instead be treated as a compiler output in its own right: it can inform downstream rule-selection mechanisms even where a binary candidate decision is not selective. This is an architectural implication, not an implemented acceleration or a claim about current Snort capabilities. It is consistent with the broader interest in coordinated programmable-data-plane and IDS processing [5].

9 Threats to Validity and Limitations

9.1 Construct Validity

The reported metrics characterize different constructs and should not be collapsed into a single notion of offloading effectiveness. Semantic coverage measures whether a safe P4 predicate was derived under the supported profile; it does not measure semantic equivalence. In particular, the 3,996 NECESSARY classifications do not constitute full Snort rule offload, and EXACT=0 means that no evaluated rule had all of its semantics represented by the P4 predicate. This is both a limitation of the evaluated target/profile and a consequence of the conservative classifier, which does not label partial preservation as equivalence.

The NECESSARY classification itself rests on a model of Snort's header evaluation (Section 4.5): which rules are service-based, which header fields Snort skips for them, and that the `flags` option matches only TCP. The model was established from the Snort 3.12.2.0 source and documentation, and confirmed on BMv2 for one service and one port: 190 HTTP rules alerted on port 18080, and under the current compilation every such witness was a candidate or belonged to a rule in fallback. An earlier compilation, which treated a service as an additional delegated condition, over-constrained 2,931 predicates, and its strong-discriminator plan marked all 190 witnesses no-match. A model that missed another way in which Snort skips header conditions would fail in the same way, and a different Snort version could change these rules; for inspector buffers whose registration was not traced, the compiler counts the rule as service-based, which only broadens its predicate.

Similarly, the 94.18% backend reduction measures P4 installation state, not traffic reduction or removal of rules from Snort. Packet-level selectivity measures candidate rates only among parser-eligible frames, whereas candidate-rule-space reduction is an offline feasibility metric and not a runtime execution reduction. Observed alert preservation is empirical and finite rather than formal equivalence. No throughput, latency, or speedup was measured. Finally, CICIDS2017 attack-window names do not by themselves establish the semantic meaning of the Snort alerts observed in those windows; an SQLi, XSS, or Brute Force window should not be interpreted as evidence that the corresponding attack class was detected unless the fired rule semantics support that conclusion.

9.2 Internal Validity

SQLi and XSS were used as development and diagnostic traces, whereas Brute Force was preserved as a held-out trace. No policy retuning was performed after inspecting the held-out trace. The strong-discriminator configuration is an operational ablation rather than a new semantic classification, and its criterion and the grouping strategies were fixed before held-out validation under the earlier compilation and applied unchanged to the current one. Experiment artifacts were frozen, with integrity checks recorded where applicable, and the earlier compilation's artifacts are archived and regenerate byte for byte.

The causal analysis initially identified broad predicates, and the residual analysis showed that coverage persisted through the entries that remained. This reduces the risk of attributing the observed behavior only to a small number of broad rules. It does not establish causal proof beyond the performed ablations.

Flowbit dependency audits were performed for both partitions of the ruleset. In the semantic partition, all 636 flowbit producer-to-consumer edges stay inside the candidate side, so running the UNSUPPORTED rules separately does not break a flowbit. The strong-discriminator partition is not closed under flowbits: 24 edges cross it, all on one flowbit, `smb.trans2`, between two port-based rules that remain P4-eligible (SIDs 3143 and 3144) and ten `netbios-ssn` rules in fallback. A deployment that ran the two sides of that partition as separate rule sets would have to keep this flowbit component together. The absence of other cross-partition dependencies does not establish safety for arbitrary Snort rulesets or all stateful mechanisms.

Semantic split preservation has limited discriminating power in the evaluated traces. The UNSUPPORTED partition contains 21 rules, and none of them produced an alert in any of the three traces. The comparison of the UNSUPPORTED side therefore contrasted an empty expected set with an empty run, and the candidate side necessarily coincided with the FULL alert set. The reported zero missing and zero additional alerts are correct, but a partition that assigned rules incorrectly would yield the same result whenever the misassigned rules are silent on the evaluated traffic. This measure establishes that removing the UNSUPPORTED rules did not perturb the observed alerts; it does not establish that the split is correct for rules that the traces never exercised. Rulesets or traces that activate UNSUPPORTED rules would be required to give this check discriminating power. The end-to-end campaign does not close this gap: the 14 rules that both compilations classify as UNSUPPORTED admit no synthesized witness, and the campaign does not run Snort on the two partitions separately.

The campaign's witnesses are synthesized from the compiler's own IR, so the generator and the compiler could share an error. The verdict does not rest on that IR, however: a case counts only when Snort, running the original rule text, alerts on it, so an IR error that produced a wrong witness yields a VACUOUS case rather than a CONFIRMED one. Vacuous cases exposed two IR defects in `content` handling, whitespace inside the pattern and negation written as `!` outside the quotes, and both were corrected without changing any P4 artifact. Both runs replay the trace of the final witness-generator revision, which preceded only the second correction; the only sampled rule it affected, SID 28982, was recorded as vacuous. The campaign's reach is also bounded by construction. Each rule is exercised at a single point of its packet space, always on a port of its own header, so the campaign cannot observe the service-based evaluation of Section 4.5; only the dedicated off-port test does. The 1,875 rules for which no witness can be synthesized, most of which carry the richer payload and state semantics, are not tested at all, and the sample is drawn from the synthesizable subset rather than from the full ruleset.

9.3 External Validity

These results should be interpreted within the experimental Snort variable profile, one Community ruleset/version, and the selected CICIDS2017 Thursday windows. The variable profile is experimental rather than an official canonical Snort configuration. Different Community versions, custom or private rulesets, `HOME_NET` definitions, service-port profiles, rule-base sizes, traffic distributions, IDS workloads, attack mixes, or encrypted-traffic distributions may change the semantic classification, candidate rate, and candidate rule set size. The high NECESSARY rate, the global candidate rate, and the observed reductions should therefore not be assumed to hold unchanged outside this configuration. Beyond the evaluated ruleset, the compiler was exercised at compile time only: it compiled the Snort 2.9.15.1 rules distributed by Ubuntu, both raw and converted by `snort2lua`, once the 48 rule lines using variables that Snort 3 does not define, and therefore also rejects, were removed. The ten header-only rules of that ruleset are the only EXACT classifications observed, and they agreed with Snort on all 340 sampled frame–rule pairs, with the P4 side evaluated by the offline matcher.

Only selected Thursday CICIDS2017 windows were evaluated for alerts. For the packet-level quantities alone, we also measured the entire Thursday capture offline (9,322,025 frames, 9,240,405 parser-eligible), which contains the three windows. The global plan still marks every eligible packet as a candidate, and the strong-discriminator plan 95.335%. The mean candidate set is 4.27 rules, a 99.42% reduction within the P4-eligible partition, or 18.29% over the full ruleset; protocol grouping reduces it by 19.68%. The windows are therefore not atypical of the day, but Snort alerts were not computed for the rest of it. The dataset is old relative to current production traffic, and its traffic mix may not reflect contemporary enterprise networks. Held-out Brute Force strengthens the evidence within this experimental setup, but it does not establish broad real-world generalization to production networks, other datasets, or other IDS workloads.

The trace-based alert evidence rests on few rules: the three windows fired one, two, and five distinct SIDs, all of them in fallback under the strong-discriminator plan. The end-to-end campaign widens rule coverage to 958 confirmed SIDs under the global plan and 134 under the strong-discriminator plan, but its traffic is synthetic, one short flow per rule, and says nothing about how often these rules fire, or which packets they fire on, in real traffic.

9.4 Implementation and Target Limitations

The prototype uses the BMv2/v1model software target rather than production ASIC hardware. It was not evaluated on production hardware, and the evaluation does not establish deployment scalability beyond the validated table state. No hardware timing, resource-utilization, line-rate, throughput, or latency measurement was conducted.

The implemented parser follows the current Ethernet/IPv4 path and does not

support VLAN encapsulation. TCP, UDP, and ICMP are the principal handled L4 protocols, but the implementation does not provide generic support for every encapsulation or protocol. Parser-ineligible frames are outside the candidate-rate denominator rather than being P4 no-matches. Moreover, the P4 program marks candidate metadata; it does not implement a complete forwarding or bypass architecture that safely removes Snort from the path for arbitrary traffic.

These parser limits are handled by exclusion rather than by silent approximation. IPv4 options, fragments, truncated headers, and parser errors are gated out of predicate evaluation and marked candidate through domain fallback, so the safety relation is stated over D and the pipeline fails open elsewhere. The cost is that selectivity is claimed only over D: a deployment whose traffic contains substantially more options, fragments, or tunneling than the evaluated traces would push a larger share of packets into unconditional fallback, which is safe but less informative. Extending the parser to cover those cases selectively, rather than conservatively, remains future work.

Some header conditions that the target could represent are still omitted, which only broadens their predicates: port ranges, although the port keys are ternary and a range decomposes into prefixes; ICMP type and code; and the protocol number given by `ip_proto`. Representing them would take at most 54, 123, and 3 rules out of the 3,278 that the strong-discriminator plan applies to every packet, so it could raise that plan's full-ruleset reduction by at most 1.3, 3.1, and 0.07 percentage points. ICMP type and code would also need new table keys.

Replay accounting reconciles frame for frame with the data-plane counters in all three traces under both plans, and for the alert packets replayed alone. An earlier held-out replay, without the longer veth queues and the 2 ms inter-frame delay, lost 87 of 1,108,658 frames (0.008%) in the transport layer, all within the candidate category; with them, every replay delivered and counted every frame.

9.5 Limitations of the Candidate-Set Analysis

The candidate rule set analysis reported mean reductions of 99.39–99.47% within the P4-eligible partition and 18.29–18.30% over the full ruleset. These values are offline candidate-rule-space feasibility measurements, and the traces exercise none of the retained predicates: every observed alert belongs to a fallback rule. The current prototype does not implement dynamic per-packet Snort rule activation, candidate-rule masks inside Snort, rule scheduling driven by P4 metadata, a candidate-set aware Snort execution engine, or throughput benchmarking of reduced rule sets.

Stateful semantics make dynamic execution a separate research problem. Even when an alert SID is present in the packet's candidate set, detection can depend on prior flow state, stream state, flowbits, temporal conditions, or state established by previous packets. The dependency audit reduces one specific class of risk, and it found one flowbit component that crosses the strong-discriminator partition; it does not prove dynamic safety for all Snort state mechanisms. The campaign measured one such mechanism directly: Snort attributed the alerts of

606 HTTP rules to an acknowledgment that does not satisfy the rule's header, and the strong-discriminator plan marked 445 of those acknowledgments no-match. These rules are in fallback under the current compilation, but under the earlier one two of them were prefiltered, so a packet-level decision that is safe for the rule's own packets can still withhold a packet that stream inspection needs. Likewise, the grouped candidate set was computed offline for four structural strategies only, and it does not show that every future grouping algorithm would fail; future evaluation should consider provenance, dependencies, overlap, and target resource constraints together.

10 Conclusion

This work examined how Snort 3 detection semantics can be partially compiled into predicates executable in a P4 data plane without obscuring semantic loss, and whether those predicates remain useful when composed at ruleset scale. The result is a semantics-aware, target-aware compilation workflow in which the semantic relation of each generated predicate is explicit and predicate provenance is retained through backend lowering.

For RQ1, all 4,017 evaluated rules were parsed; 3,996 were classified NECESSARY and 21 remained UNSUPPORTED, while no rule was EXACT. Under the implementation assumptions of Section 4, the construction gives $r \Rightarrow p_r$ inside the enforced packet domain, with options, fragments, and truncated headers retained by fail-open fallback, and without presenting partial preservation as equivalence. For the 2,931 service-based rules, this required compiling only their header addresses, because Snort 3 skips header ports and protocol once it identifies the flow's service. An earlier compilation that kept them over-constrained those predicates, and on BMv2 its strong-discriminator plan marked no-match the witnesses of 190 HTTP rules on which Snort alerted off their header port. For RQ2, 4,040 source predicate instances were lowered to 235 final installation entries, a 94.18% installation-state reduction, and the 235-entry plan was installed and read back on the evaluated BMv2/v1model target without encoding or validation failures.

RQ3 showed that semantic coverage and operational selectivity are distinct. Under the evaluated configuration, the binary global prefilter marked all parser-eligible packets as candidates in the SQLi, XSS, and held-out Brute Force traces, replayed on BMv2 with exact accounting; the strong-discriminator ablation still marked 92.2–94.9%. Broad necessary predicates contributed initially, and a few structurally specific entries that coincide with the dominant traffic sustained the coverage after they were removed.

For RQ4, the evaluated predicates achieved 100% observed alert preservation in the three traces, and the end-to-end campaign observed no violation of $r \Rightarrow$ cand among the sampled rules that Snort confirmed: 958 under the global plan and 134 under the strong-discriminator plan, 36 of the latter attributable to the rule's own entries alone. The same campaign showed that the global plan's broad entries leave no confirmation attributable to a single rule, and that

HTTP alerts often fall on an acknowledgment that does not satisfy the rule's header. Provenance exposed per-packet candidate rule sets of about four of the 739 P4-eligible rules, but the 3,278 fallback rules apply to every packet, and the reduction over the full ruleset is 18.29–18.30%. Under the earlier compilation, the same derivation reported 81.81% on held-out Brute Force, most of it from header ports that Snort 3 does not check. This candidate-set result is an offline feasibility result, not a dynamic IDS execution result. The central implication is that semantic correctness and operational usefulness are distinct dimensions of IDS rule offloading, and that a prefilter's apparent selectivity has to be checked against how the IDS evaluates each rule. Future work should implement candidate-set-aware and dependency-aware IDS execution, investigate selectors that use the identified service as well as header provenance, then evaluate execution cost on hardware targets and broader rulesets, profiles, traffic datasets, and parser/protocol/ encapsulation coverage.

Artifact Availability

The compiler, the experimental variable profile, the analysis scripts, and the frozen output artifacts reported in this paper are available at `https://github.com/medeiros08/Snort-p4`. The repository contains the compilation pipeline from rule parsing to P4Runtime entry generation, the semantic manifest with predicate-to-rule provenance, the compiled P4 program and its P4Info, and the JSON artifacts underlying every table and figure. A regression check reproduces the ten pipeline artifacts byte for byte from the committed sources; the earlier compilation's artifacts are archived and regenerate byte for byte from the same sources with its header semantics selected. The end-to-end campaign's cases, per-case verdicts, and Snort alerts are under `output/campaign/`; its traces are not versioned, but the case generator at the commit that produced a run should regenerate its trace byte for byte; this was checked against the SHA-256 digest recorded with the run.

The CICIDS2017 traces are not redistributed here; they are obtained from the Canadian Institute for Cybersecurity. The BMv2 target, Snort 3, and the P4 compiler are third-party software and are not redistributed; the exact versions used to produce the frozen artifacts are stated in Section 6, and `scripts/setup_bmv2.sh` reproduces the data-plane configuration they ran under.

A Experimental Snort Variable Profile

The evaluated CICIDS2017 configuration used the `cicids2017` profile from `config/profiles/cicids2017.json`. Its resolved variables were written to `config/snort_variables.json`. The generated Snort configuration is `config/snort3_experiment.lua`; the active-profile record identifies the same profile. Table 11 documents the resulting variables used for the 4,017-rule evaluation.

The generated configuration includes the standard Snort configuration and

Table 11: Experimental Snort variable profile used for the CICIDS2017 evaluation.

Variable	Experimental value	Role
HOME_NET	192.168.10.0/24	protected network
EXTERNAL_NET	any	external network
HTTP_SERVERS	192.168.10.0/24	HTTP server scope
SQL_SERVERS	192.168.10.0/24	SQL server scope
SMTP_SERVERS	192.168.10.0/24	SMTP server scope
TELNET_SERVERS	192.168.10.0/24	TELNET server scope
SIP_SERVERS	192.168.10.0/24	SIP server scope
HTTP_PORTS	80 8080	HTTP service ports
ORACLE_PORTS	1521	Oracle service ports
FILE_DATA_PORTS	80 110 143	file-data service ports
SIP_PORTS	5060 5061	SIP service ports
SSH_PORTS	22	SSH service ports
FTP_PORTS	21	FTP service ports

then explicitly disables the default inclusion of `file_magic.rules` by setting `file_inspect = {}`. Consequently, the evaluated ruleset contains only the intended 4,017 community rules.

This is an experimental variable profile, not an official canonical Snort configuration. The semantic classifications and selectivity measurements are configuration dependent: changing network variables or service-port definitions can change the generated necessary predicates and their selectivity. This profile-specific scope is considered in the threats to validity discussion.

References

[1] Ali AlSabeh, Joseph Khoury, Elie Kfoury, Jorge Crichigno, and Elias Bou-Harb. A survey on security applications of P4 programmable switches and a STRIDE-based vulnerability assessment. *Computer Networks*, 207:108800, 2022.

[2] Pat Bosshart, Dan Daly, Glen Gibb, Martin Izzard, Nick McKeown, Jennifer Rexford, Cole Schlesinger, Dan Talayco, Amin Vahdat, George Varghese, and David Walker. P4: Programming protocol-independent packet processors. *ACM SIGCOMM Computer Communication Review*, 44(3):87–95, 2014.

[3] H. B. Brum, L. A. D. Knob, T. Ferreto, and D. Siracusa. Rethinking NIDS rule-based pre-filtering. In *IFIP/IEEE International Symposium on Integrated Network Management and Network Operations (CNSM)*, 2025.

[4] Mah-Rukh Fida and Azza H. Ahmed. Leveraging programmable data plane for network intrusion detection: A survey. *Computer Networks*, 281:112202, 2026.

[5] L. S. Hagen, A. S. Ilha, R. Parizotto, and L. P. Gaspary. RNA: Automating IDS/IPS event detection offload into programmable forwarding devices. In *IEEE/IFIP Network Operations and Management Symposium (NOMS)*, 2025.

[6] Frederik Hauser, Marco Häberle, Daniel Merling, Steffen Lindner, Vladimir Gurevich, Florian Zeiger, Reinhard Frank, and Michael Menth. A survey on data plane programming with P4: Fundamentals, advances, and applied research. *Journal of Network and Computer Applications*, 212:103561, 2023.

[7] Benjamin Lewis, Matthew Broadbent, and Nicholas Race. P4ID: P4 enhanced intrusion detection. In *2019 IEEE Conference on Network Function Virtualization and Software Defined Networks (NFV-SDN)*, 2019.

[8] Benjamin Lewis, Matthew Broadbent, Charalampos Rotsos, and Nicholas Race. 4MIDable: Flexible network offloading for security VNFs. *Journal of Network and Systems Management*, 31:52, 2023.

[9] Ali Mazloum, Ali AlSabeh, Elie Kfoury, and Jorge Crichigno. Security applications in P4: Implementation and lessons learned. *Computer Networks*, 257:111011, 2025.

[10] p4lang. BMv2: The behavioral model software switch. `https://github.com/p4lang/behavioral-model`, 2025. `simple_switch_grpc` 1.15.0, commit 68f4a978.

[11] Martin Roesch. Snort: Lightweight intrusion detection for networks. In *Proceedings of the 13th USENIX Conference on System Administration (LISA '99)*, pages 229–238, Seattle, WA, USA, 1999. USENIX Association.

[12] Eduardo Vudala Senoski, João Ribeiro Andreotii, Raphael Kaviak Machnicki, Jorge Pires Correia, and Vinicius Fulber-Garcia. P4NIPS: P4 network intrusion prevention system. In *Anais Estendidos do Simpósio Brasileiro de Cibersegurança (SBSeg)*, pages 75–84, 2025.

[13] Iman Sharafaldin, Arash Habibi Lashkari, and Ali A. Ghorbani. Toward generating a new intrusion detection dataset and intrusion traffic characterization. In *Proceedings of the 4th International Conference on Information Systems Security and Privacy (ICISSP)*, pages 108–116, 2018.

[14] Kairo Tavares and Tiago Ferreto. P4-ONIDS: A P4-based NIDS optimized for constrained programmable data planes in SDN. In *Anais do XXXIX Simpósio Brasileiro de Redes de Computadores e Sistemas Distribuídos*, pages 434–447, 2021.

[15] The P4 Language Consortium. *$P4_{16}$ Language Specification*, 2023. Version 1.2.4, `https://p4.org/specs/`.

[16] The P4.org API Working Group. *P4Runtime Specification*, 2024. Version 1.4.1, `https://p4.org/specs/`.

[17] The Snort Team. *Snort 3 User Manual*, 2026. Snort 3.12.2.0, `https://github.com/snort3/snort3`.